

# OPENCL编程模型

---

汤善江 副教授

天津大学智能与计算学部

[tashj@tju.edu.cn](mailto:tashj@tju.edu.cn)

<http://cic.tju.edu.cn/faculty/tangshanjiang/>

# Outline

- OpenCL概述
- OpenCL抽象模型
  - 平台模型
  - 执行模型
  - 内存模型
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例

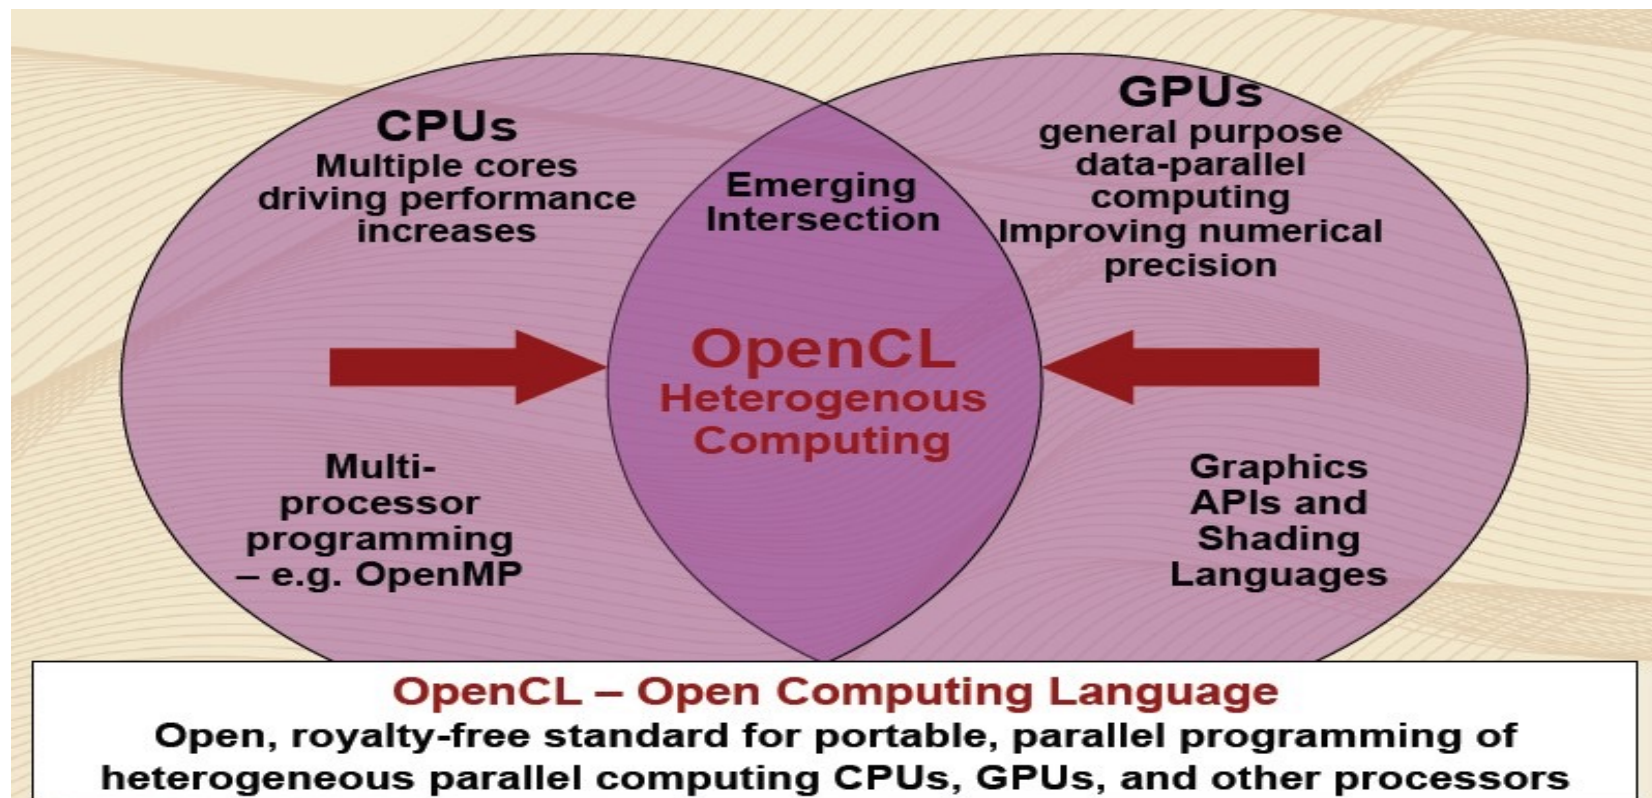
# Outline

- **OpenCL概述**
- OpenCL抽象模型
  - 平台模型
  - 执行模型
  - 内存模型
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例

# 异构计算系统

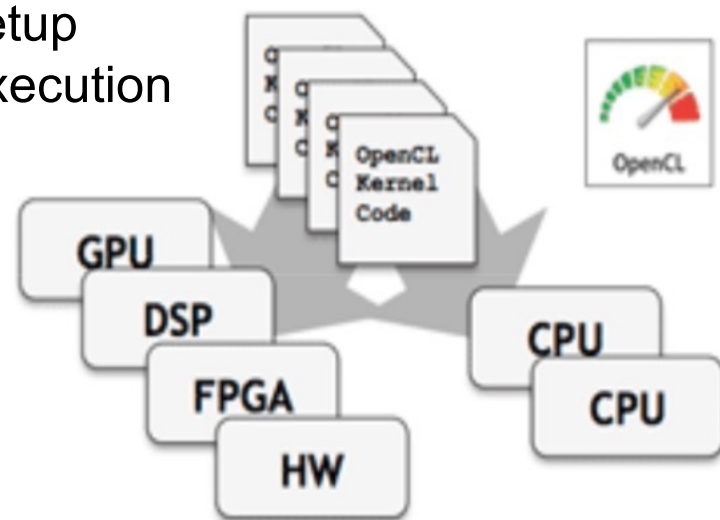
- A modern computer system has:
  - CPU(s)
  - GPU(s)
  - DSP processors
  - ... other?
- Need to make the best use of **all** the available resources from within a **single** program:
  - One program that runs well (i.e., reasonably close to “hand-tuned” performance) on a heterogeneous mixture of processors.

# 异构计算系统编程



# 什么是OpenCL?

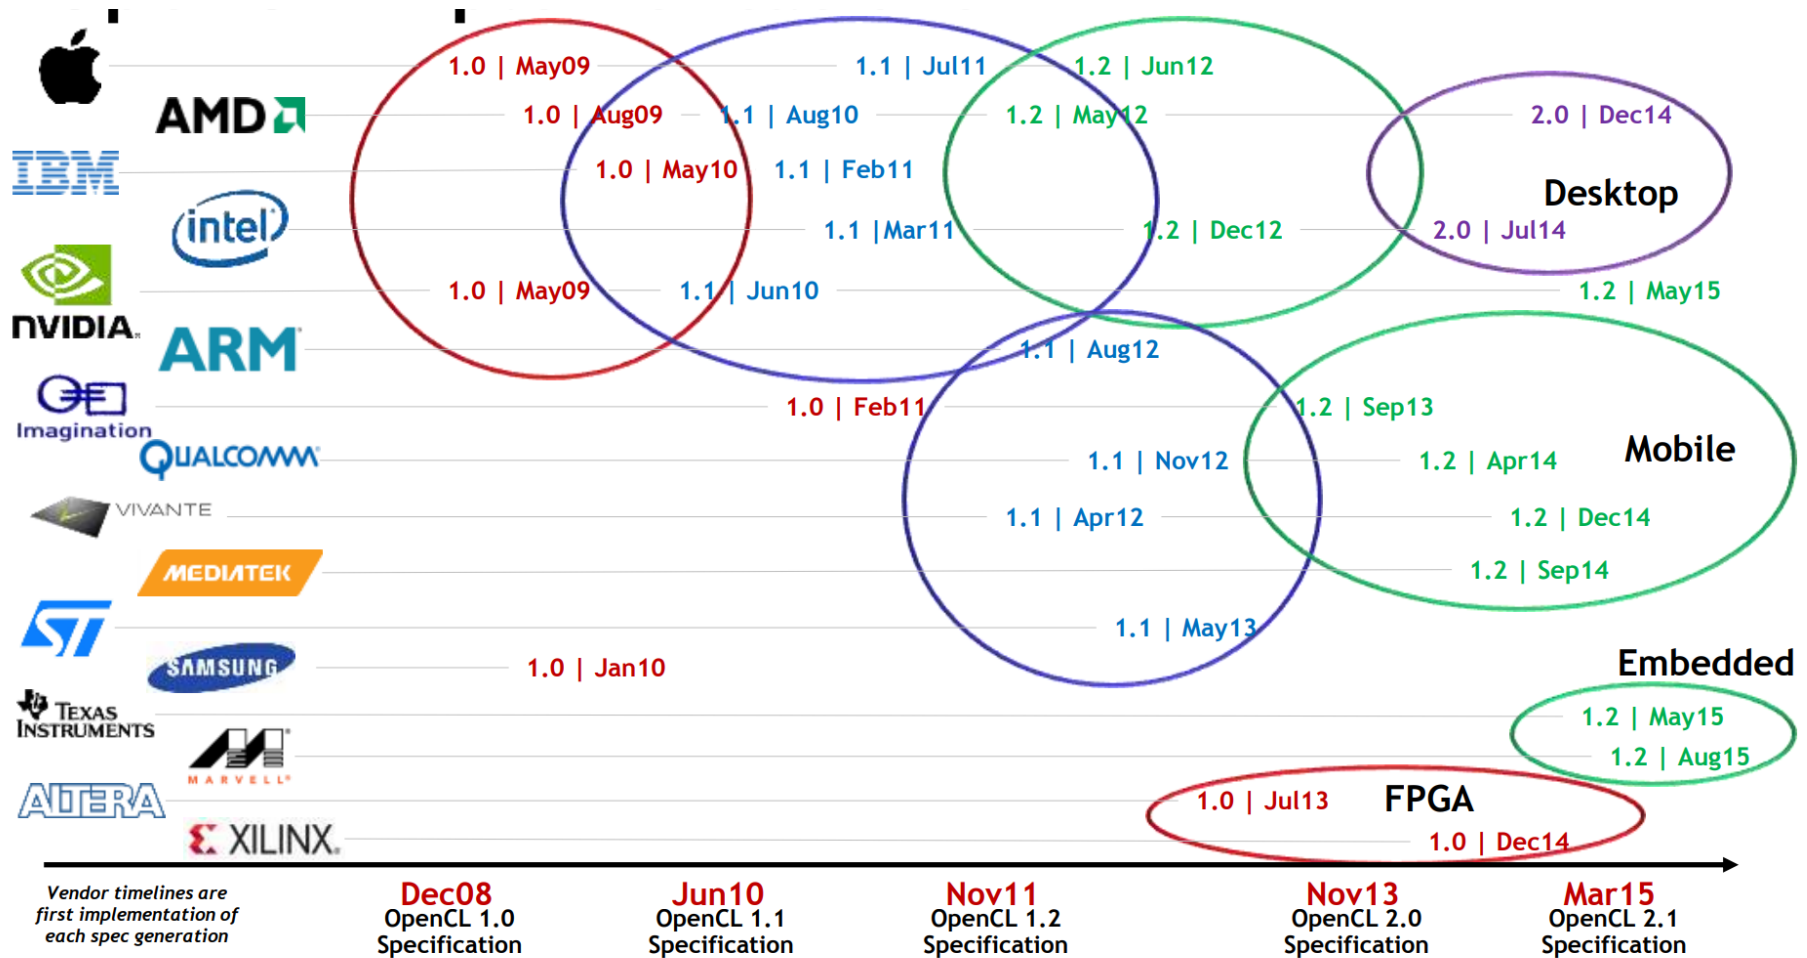
- OpenCL (Open Computing Language) 是一个为异构平台编写程序的框架，此异构平台可由CPU，GPU或其他类型的处理器组成。
- Low-level programming API for data parallel computation
  - Platform API: device query and pipeline setup
  - Runtime API: resources management + execution
- **Cross-platform** API
  - Windows, MAC, Linux, Mobile, Web...
- **Portable** device targets
  - CPUs, GPUs, FPGAs, DSPs, etc...
  - **One code** tree can be executed on CPUs, GPUs, DSPs, FPGA and hardware
- Implementation based on C99
- Maintained by Khronos Group ([www.khronos.org](http://www.khronos.org))
- Current version: 3.0 with C++ support (classes & templates)



# OpenCL的发展历史

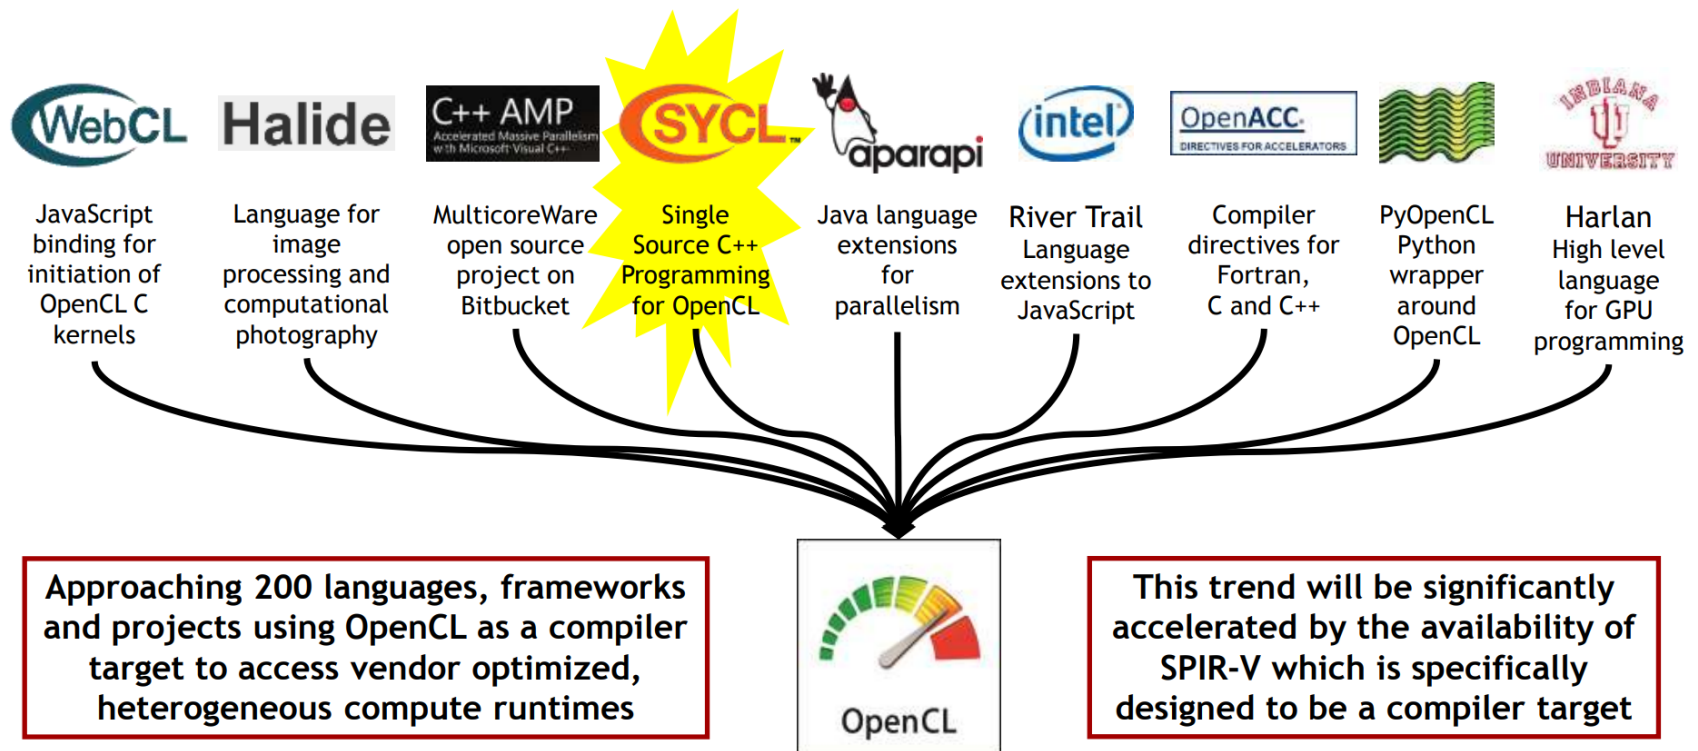
- 2008年 Apple  
提出OpenCl规范
- 2008年 Khronos Group  
发布OpenCL 1.0标准
- 2013年 Khronos Group  
OpenCl 2.0标准发布
- 2013年 截至2013年11月，以AMD、Apple、ATI、Intel、Nvidia等为代表的公司，已经发布了多款支持OpenCl的产品。
- 2017年 Khronos Group  
OpenCl 2.2标准发布
- 2020年 Khronos Group  
OpenCl 3.0标准发布

# OpenCL Implementations





# OpenCL Front-End APIs



# OpenCL: A high-level view

- OpenCL applications:
  - A host program running on the PC
  - One or more Kernels that are queued up to run on CPUs, GPUs, and “other processors”.
- OpenCL is understood in terms of these models
  - Platform model
  - Execution model
  - Memory model
  - Programming model

# Outline

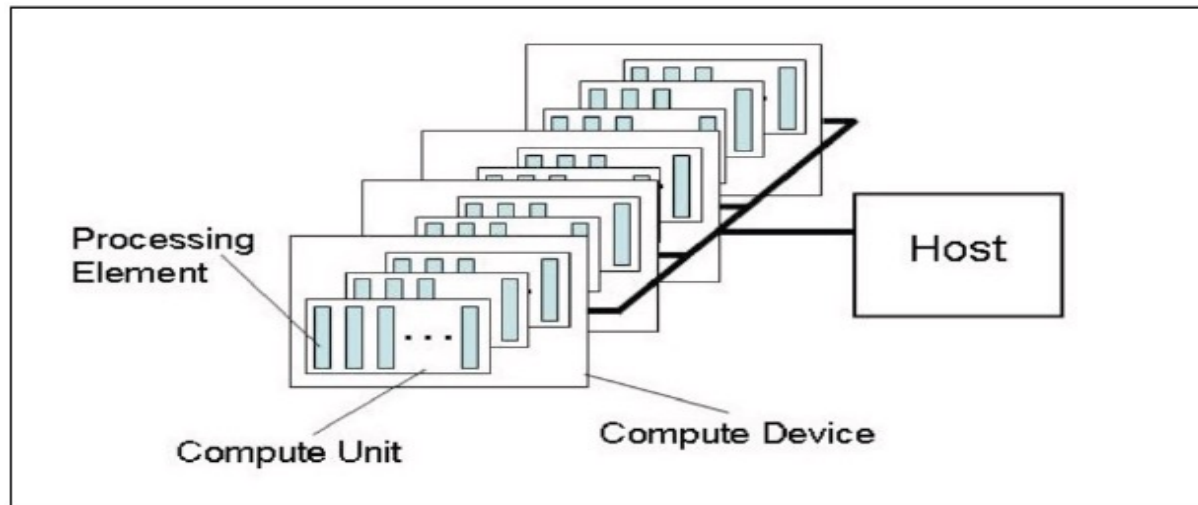
- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - 内存模型
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例

# OpenCL 平台模型

- Each OpenCL implementation (i.e. an OpenCL library from AMD, NVIDIA, etc.) defines *platforms* which enable the host system to interact with OpenCL-capable devices
  - Currently each vendor supplies only a single platform per implementation
- OpenCL uses an “Installable Client Driver” model
  - The goal is to allow platforms from different vendors to co-exist
  - Current systems’ device driver model will not allow different vendors’ GPUs to run at the same time

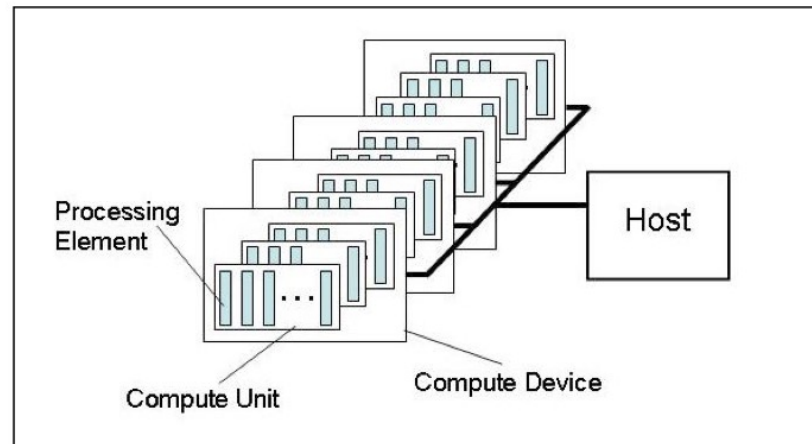
# OpenCL 平台模型

- One Host + one or more Compute Devices
  - Each Compute Device is composed of one or more Compute Units
    - Each Compute Unit is further divided into one or more Processing Elements
      - PE executes code as SIMD or SPMD

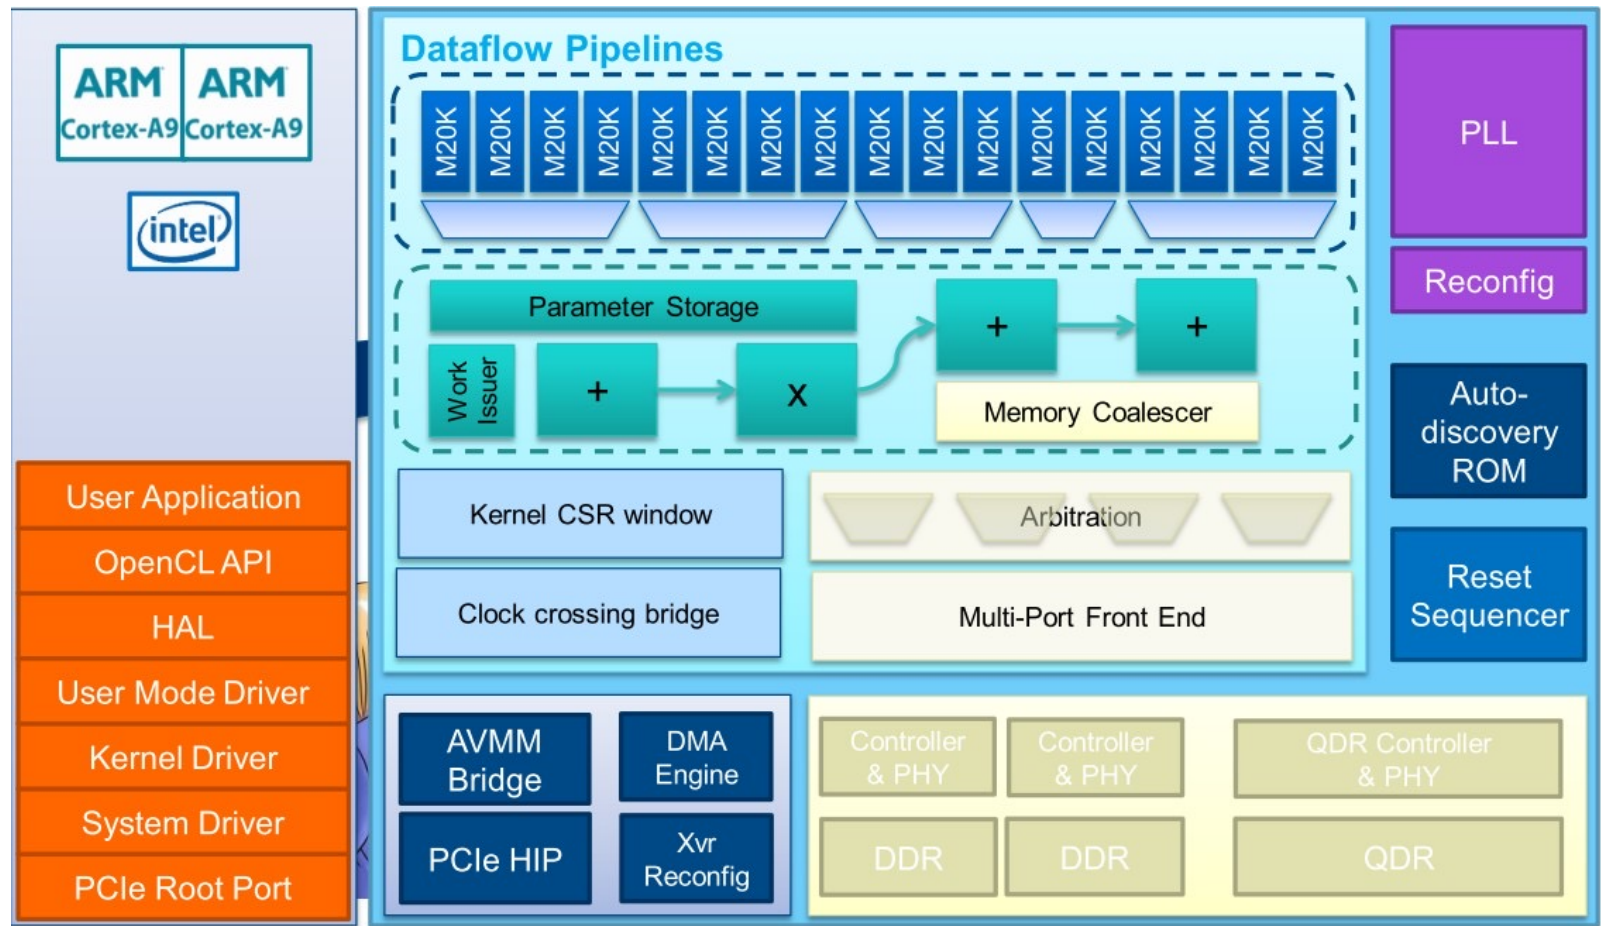


# OpenCL 平台模型

- The host is whatever the OpenCL library runs on
  - x86 CPUs for both NVIDIA and AMD
- Devices are processors that the library can talk to
  - CPUs, GPUs, and generic accelerators
- For AMD
  - All CPUs are combined into a single device (each core is a compute unit and processing element)
  - Each GPU is a separate device



# Altera OpenCL 平台模型



Developers mostly working happily at Application level in C for kernels and favourite application programming language

The OpenCL compiler takes care of much of the worries including machine-based profiling, tuning and optimization

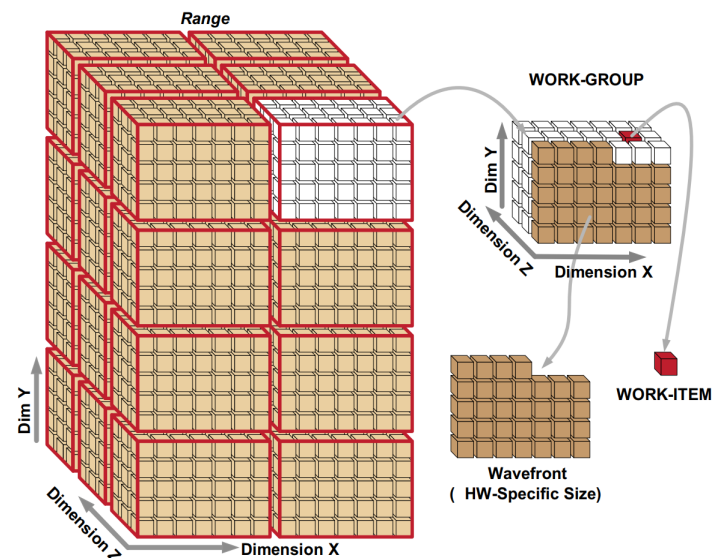
# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - **执行模型**
  - 内存模型
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例



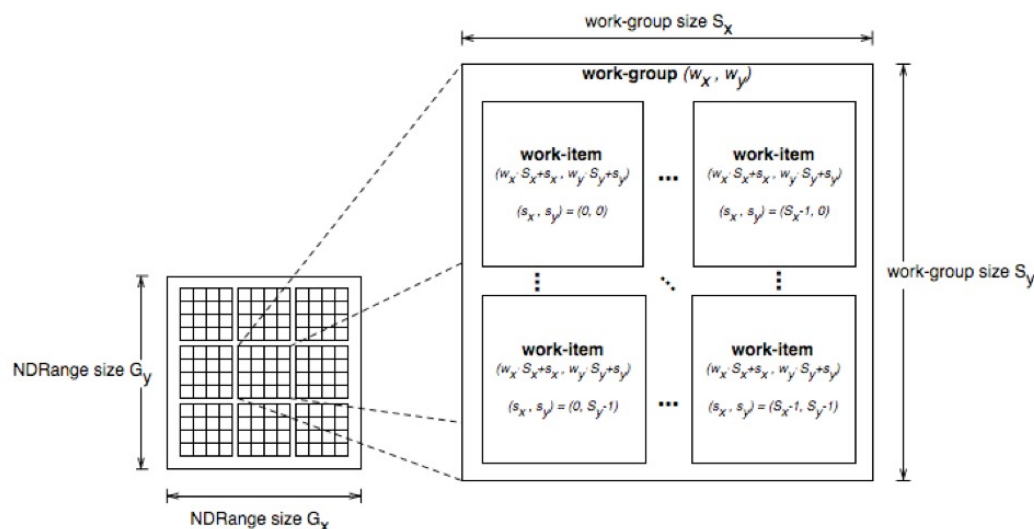
# OpenCL 执行模型

- A **kernel** is logical unit of instructions to be executed on a compute device.
- Kernels are executed in multi-dimensional **index space**:  
**NDRange**
- For every element of the index space a **work-item** is executed
- The index space is tiled into **work-groups**
- Work items within a workgroup are synchronized using barriers or fences



# OpenCL 执行模型

- Host defines a **command queue** and associates it with a context (devices, kernels, memory, etc).
- Host enqueues commands to the command queue
- Work-items can uniquely identify themselves based on:
  - A global id (unique within the index space)
  - A work-group ID and a local ID within the work-group



# OpenCL 线程组织结构

- Threads can determine their global ID in each dimension
  - `get_global_id(dim)`
  - `get_global_size(dim)`
- Or they can determine their work-group ID and ID within the workgroup
  - `get_group_id(dim)`
  - `get_num_groups(dim)`
  - `get_local_id(dim)`
  - `get_local_size(dim)`
- `get_global_id(0) = column, get_global_id(1) = row`
- `get_num_groups(0) * get_local_size(0) == get_global_size(0)`

# OpenCL执行模型——与CUDA映射

- OpenCL terminology aims for generality

OpenCL Terminology	CUDA Terminology
Compute Unit	Streaming Processor (SM)
Processing Element	Processor Core
Kernel	Kernel
Host program	Host program
Wavefront (AMD)	Warp
Work-item	Thread
Work-group	Thread Block
NDRange	Grid

# OpenCL执行模型——与CUDA映射

- Work Items Indexing

OpenCL Terminology	CUDA Terminology
get_num_groups()	gridDim
get_local_size()	blockDim
get_group_id()	blockIdx
get_local_id()	threadIdx
get_global_id()	$\text{blockIdx} * \text{blockDim} + \text{threadIdx}$
get_global_size()	$\text{gridDim} * \text{blockDim}$

# OpenCL执行模型——与CUDA映射

- Threads Synchronization

OpenCL Terminology	CUDA Terminology
barrier()	__syncthreads()
No direct equivalent*	__threadfence()
mem_fence()	__threadfence_block()
No direct equivalent*	__threadfence_system()
No direct equivalent*	__syncwarp()
Read_mem_fence()	No direct equivalent*
Write_mem_fence()	No direct equivalent*

# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - **内存模型**
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例

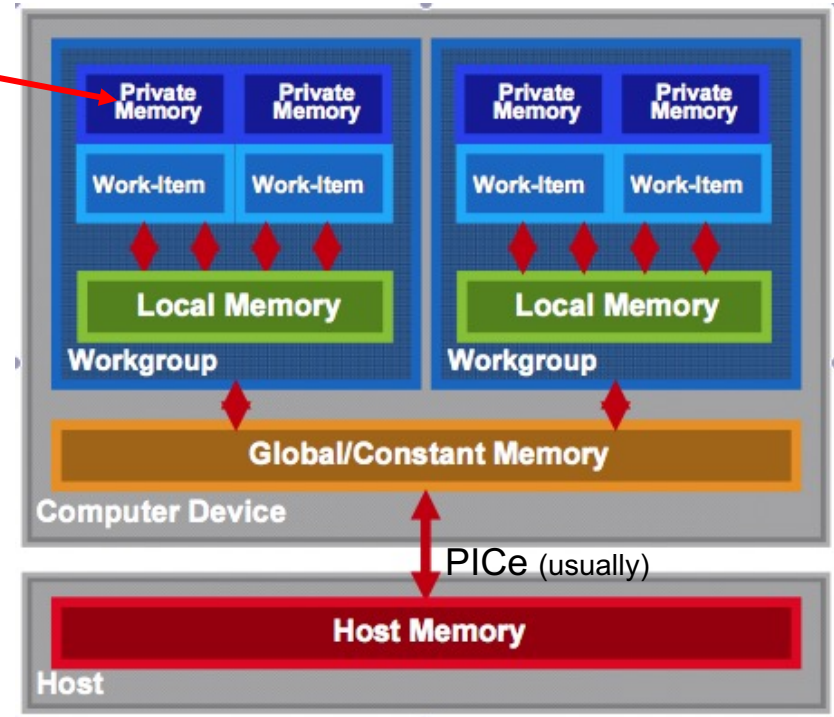
# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - **内存模型**
  - 编程模型
- OpenCL程序设计
  - 程序设计流程
  - 编程实例



# OpenCL 内存模型

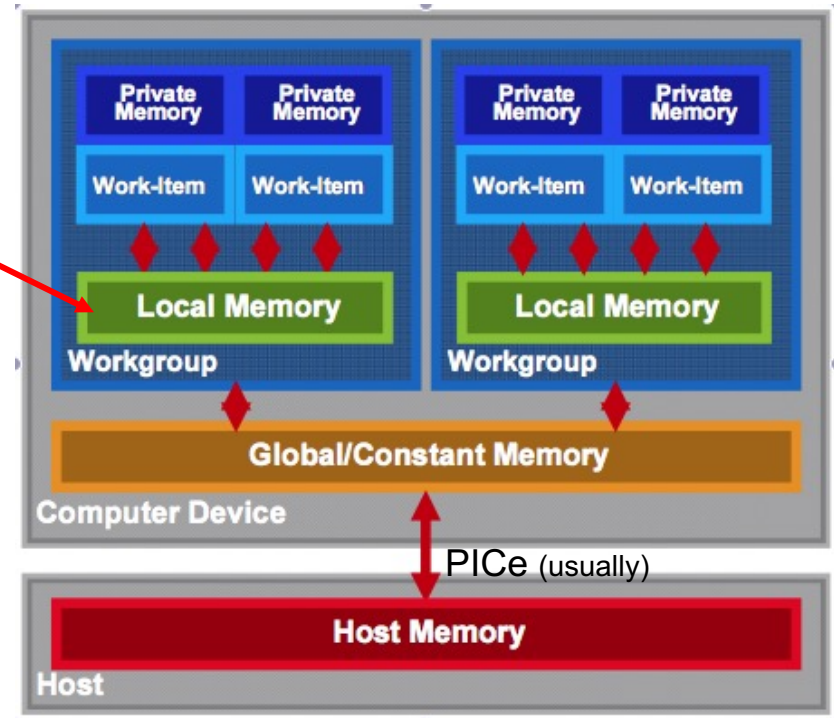
- **Private memory:** available per work item
- Local memory: shared in workgroup
- NB: No synchronization between workgroups
- Synchronization possible between work items in a common workgroup
- Global/constant memory accessible by any work-items (no guarantee to be synchronized)
- Host memory: access through the CPU



- Memory management is explicit...
- Data moved from:  
host->global->local and back

# OpenCL 内存模型

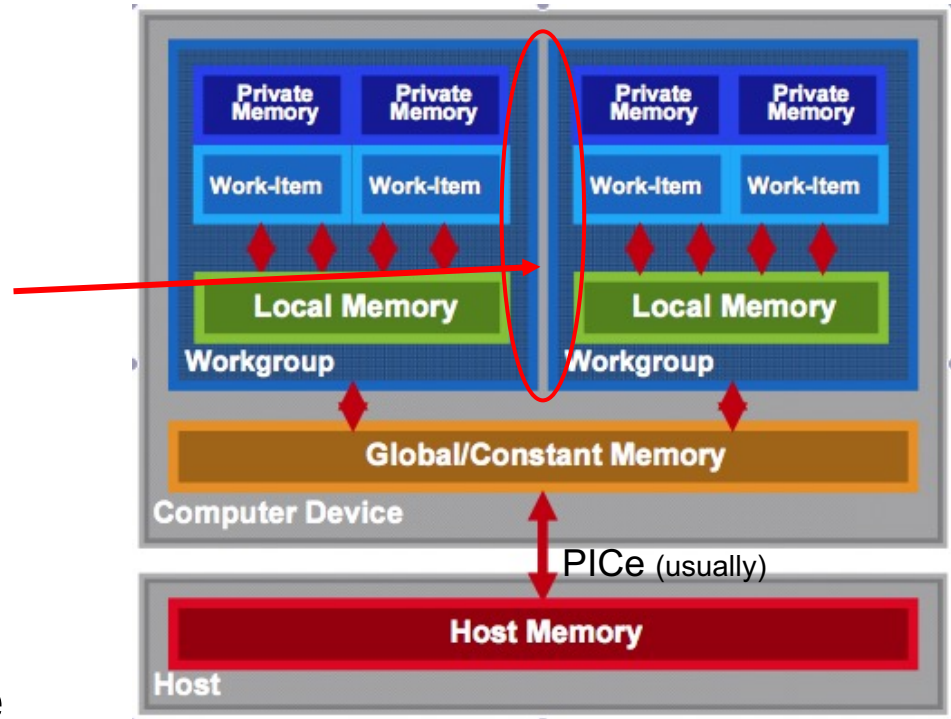
- Private memory: available per work item
- **Local memory:** shared in workgroup
- NB: No synchronization between workgroups
- Synchronization possible between work items in a common workgroup
- Global/constant memory accessible by any work-items (no guarantee to be synchronized)
- Host memory: access through the CPU



- Memory management is explicit...
- Data moved from: host->global->local and back

# OpenCL 内存模型

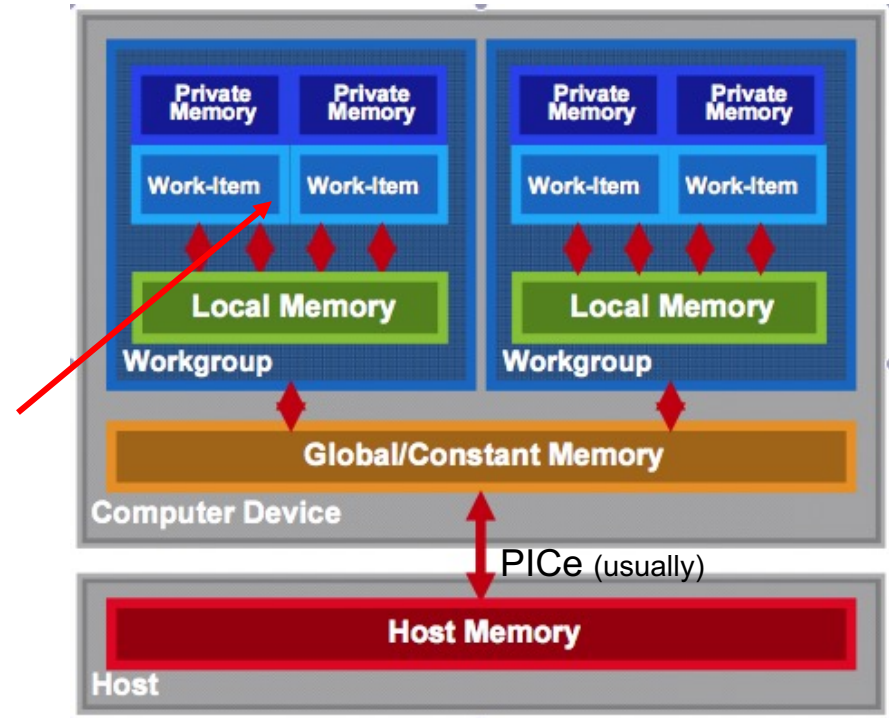
- Private memory: available per work item
- **Local memory:** shared in workgroup
- **NB: No synchronization** between workgroups
- Synchronization possible between work items in a common workgroup
- Global/constant memory accessible by any work-items (no guarantee to be synchronized)
- Host memory: access through the CPU



- Memory management is explicit...
- Data moved from: host->global->local and back

# OpenCL 内存模型

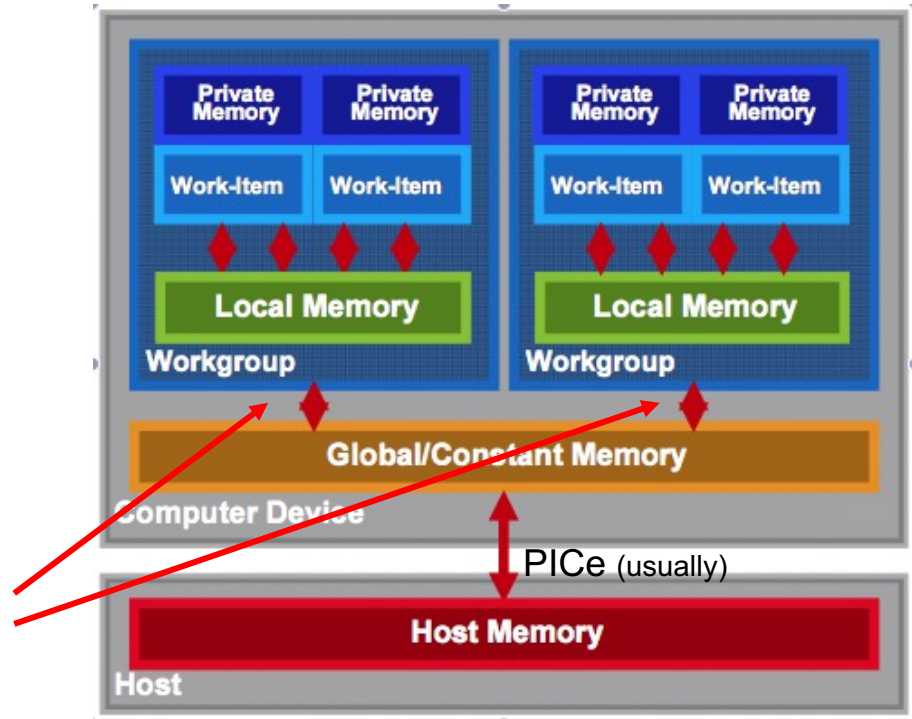
- Private memory: available per work item
- Local memory: shared in workgroup
- NB: No synchronization between workgroups
- **Synchronization possible between work items in a common workgroup**
- Global/constant memory accessible by any work-items (no guarantee to be synchronized)
- Host memory: access through the CPU



- Memory management is explicit...
- Data moved from: host->global->local and back

# OpenCL 内存模型

- Private memory: available per work item
- Local memory: shared in workgroup
- NB: No synchronization between workgroups
- Synchronization possible between work items in a common workgroup
- **Global/constant memory accessible by any work-items (no guarantee to be synchronized)**
- Host memory: access through the CPU

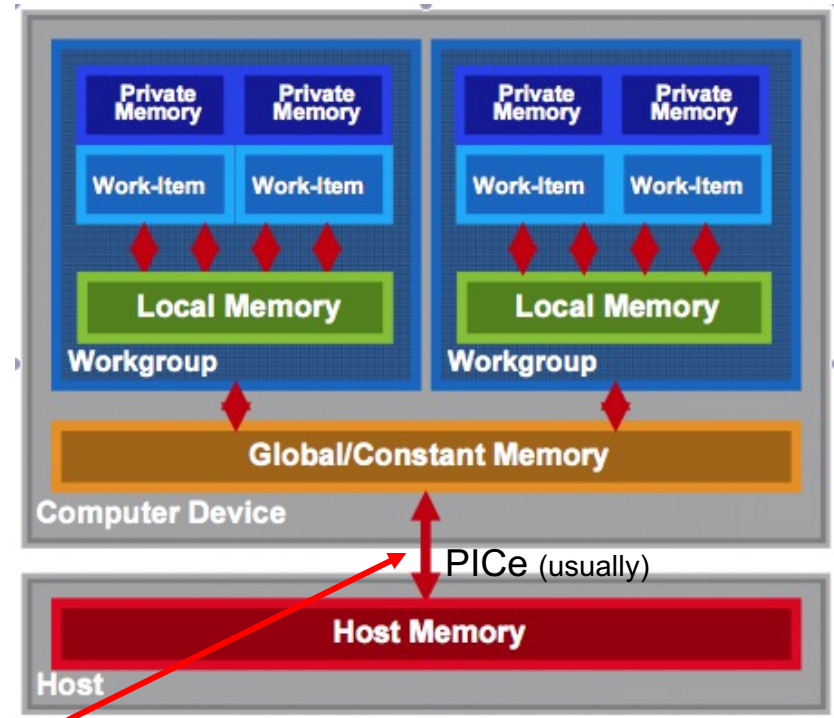


- Memory management is explicit...
- Data moved from: host->global->local and back



# OpenCL 内存模型

- Private memory: available per work item
- Local memory: shared in workgroup
- NB: No synchronization between workgroups
- Synchronization possible between work items in a common workgroup
- Global/constant memory accessible by any work-items (no guarantee to be synchronized)
- **Host memory: access through the CPU**



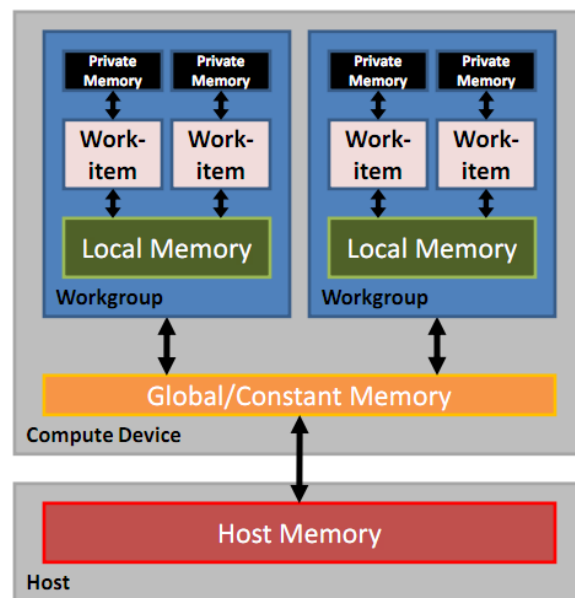
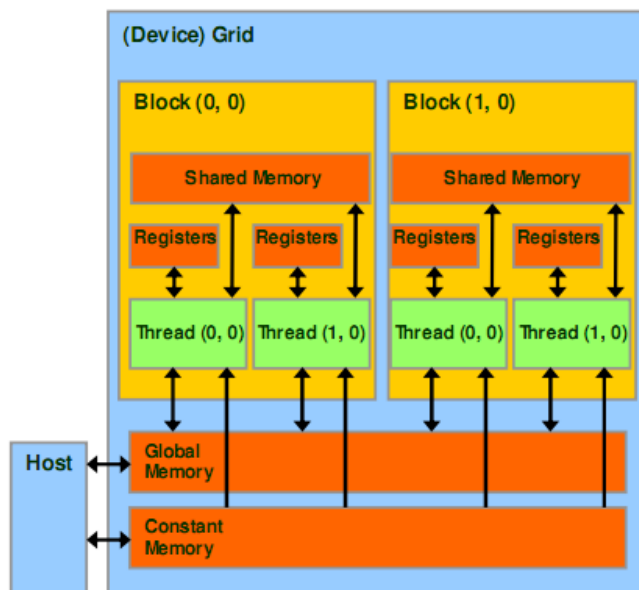
- Memory management is explicit...
- Data moved from: host->global->local and back

# OpenCL 内存模型

- Memory management is explicit
  - Must move data from host memory to device global memory, from global memory to local memory, and back
- Work-groups are assigned to execute on compute-units
  - No guaranteed communication/coherency between different work-groups (no software mechanism in the OpenCL specification)

# OpenCL 内存模型——与CUDA映射

OpenCL Terminology	CUDA Terminology
Global Memory	Global Memory
Constant Memory	Constant Memory
Local Memory	Shared Memory
Private Memory	Local Memory





# OpenCL 内存模型——与CUDA映射

- Resources Qualifiers

Description	OpenCL Terminology	CUDA Terminology
Kernel global function	<code>__kernel</code>	<code>__global__</code>
Kernel local function	nothing*	<code>__device__</code>
Readonly memory	<code>__constant</code>	<code>__device__</code>
Global memory	<code>__global</code>	<code>__device__</code>
Private memory	<code>__local</code>	<code>__shared__</code>

# CUDA与OpenCL关键指标对应汇总

## CUDA

- Thread
- Thread-block
- Global memory
- Constant memory
- Shared memory
- Local memory
- `__global__` function
- `__device__` function
- `__constant__` variable
- `__device__` variable
- `__shared__` variable



## OpenCL

- Work-item
- Work-group
- Global memory
- Constant memory
- Local memory
- Private memory
- `__kernel` function
- no qualification needed
- `__constant` variable
- `__global` variable
- `__local` variable

# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - 内存模型
  - **编程模型**
- OpenCL程序设计
  - 程序设计流程
  - 编程实例

# OpenCL编程模型

- 数据并行化
  - Work-items in a work-group run the same program
  - Update data structures in parallel using the work-item ID to select data and guide execution.
- 任务并行化
  - One work-item per work group ... for coarse grained task-level parallelism.
  - Native function interface: trap-door to run arbitrary code from an OpenCL command-queue.
  - Multiple different kernels can be executed in parallel

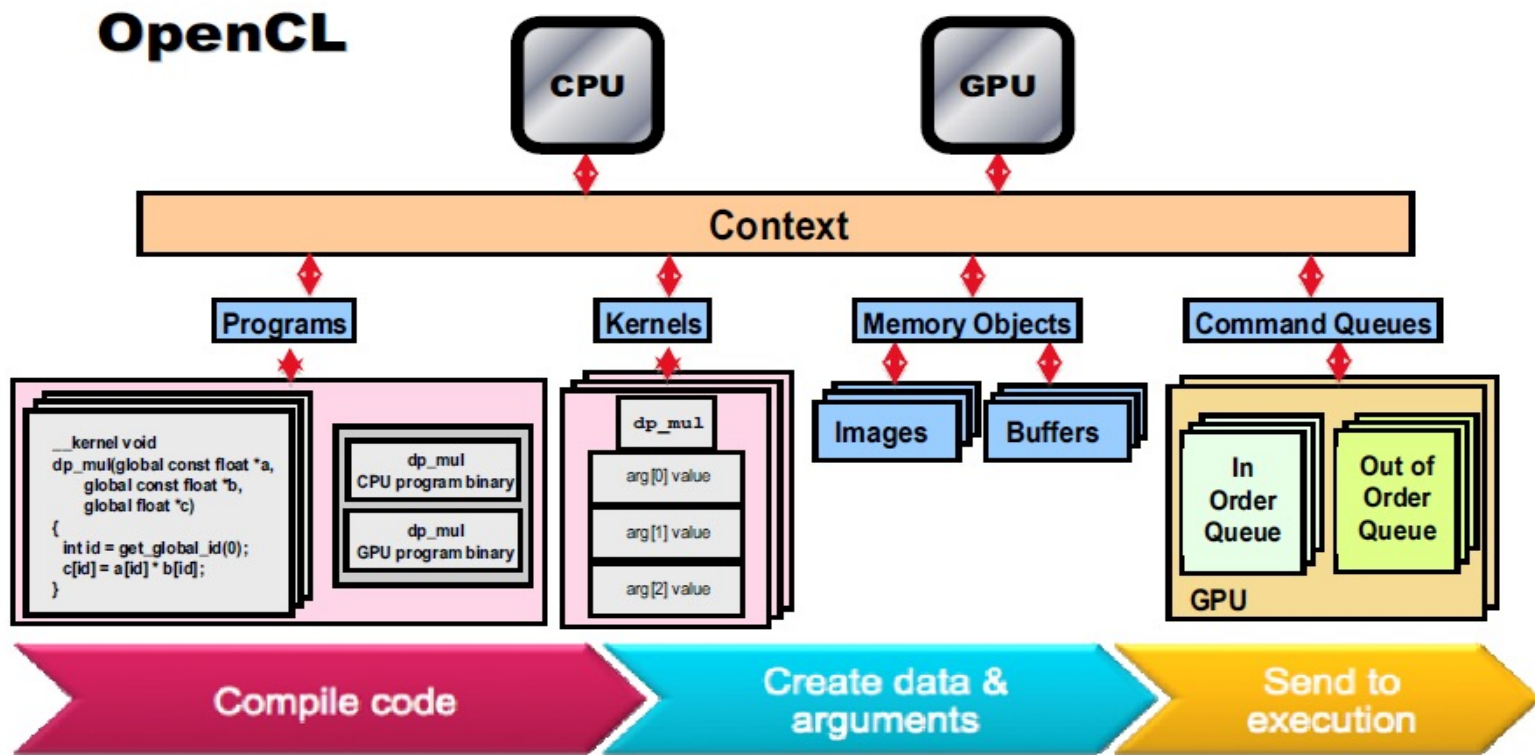
# OpenCL 数据并行化

- Kernels executed across a global domain of **work-items**
  - **Global dimensions** define the range of computation
  - One work-item per computation, executed in parallel
  - Work-items are grouped in local **work-groups**
    - **Local dimensions** define the size of the work-groups
    - Executed together on one device
    - Share local memory and synchronization
  - Caveats
    - Global work-items must be independent: **No global synchronization**
    - Synchronization can be done within a work-group

# OpenCL 任务并行化

- Expressing Task-Parallelism in OpenCL
  - Execute as a single work-item
  - A kernel written in OpenCL C
- `clEnqueueTask`
  - Imagine “sea of different tasks” executing concurrently
  - A task “owns the core” (i.e., a workgroup size of 1)

# OpenCL Program Flow



# OpenCL vs CUDA

- API Terminology

OpenCL Terminology	CUDA Terminology
<code>clGetContextInfo()</code>	<code>cuDeviceGet()</code>
<code>clCreateCommandQueue()</code>	No direct equivalent*
<code>clBuildProgram()</code>	No direct equivalent*
<code>clCreateKernel()</code>	No direct equivalent*
<code>clCreateBuffer()</code>	<code>cuMemAlloc()</code>
<code>clEnqueueWriteBuffer()</code>	<code>cuMemcpyHtoD()</code>
<code>clEnqueueReadBuffer()</code>	<code>cuMemcpyDtoH()</code>
<code>clSetKernelArg()</code>	No direct equivalent*
<code>clEnqueueNDRangeKernel()</code>	<code>kernel&lt;&lt;&lt;...&gt;&gt;&gt;()</code>
<code>clReleaseMemObj()</code>	<code>cuMemFree()</code>



# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - 内存模型
  - 编程模型
- **OpenCL程序设计**
  - 程序设计流程
  - 编程实例

# OpenCL程序设计流程

1 查询平台	<code>clGetPlatformIDs()</code>
2 查询设备	<code>clGetDeviceIDs()</code>
3 创建上下文：将平台设备与上下文关联起来	<code>clCreateContext()</code> 或 <code>clCreateContextFromType()</code>
4 创建命令队列	<code>clCreateCommandQueue()</code>
5 读取、编译内核	<code>clCreateProgramWithSource()</code> 或 <code>clCreateProgramWithBinary()</code> <code>clBuildProgram()</code>
6 打包生成内核	<code>clCreateKernel()</code>
7 创建缓存对象或图像对象， 为内核参数分配内存	<code>clCreateBuffer()</code> 、 <code>clCreateSubBuffer()</code> <code>clCreateImage2D()</code> 、 <code>clCreateImage3D()</code>
8 设置内核参数，将上面分配的内存发送到设备上	<code>clSetKernelArg()</code>
9 执行内核	<code>clEnqueueTask()</code> 或 <code>clEnqueueNDRangeKernel()</code>
10 读取设备上的处理结果	<code>clEnqueueReadBuffer()</code>
11 释放创建的资源： 创建的内存、 命令队列、 内核、 打包的程序、 上下文	<code>clReleaseMemObject()</code> 、 <code>clReleaseCommandQueue()</code> 、 <code>clReleaseKernel()</code> 、 <code>clReleaseProgram()</code> 、 <code>clReleaseContext()</code>

# Selecting a Platform

```
cl_int  clGetPlatformIDs (cl_uint num_entries,  
                           cl_platform_id *platforms,  
                           cl_uint *num_platforms)
```

- This function is usually called twice
  - The first call is used to get the number of platforms available to the implementation
  - Space is then allocated for the platform objects
  - The second call is used to retrieve the platform objects

# Selecting Devices

- Once a platform is selected, we can then query for the devices that it knows how to interact with

```
clGetDeviceIDs4 (cl_platform_id platform,  
                 cl_device_type device_type,  
                 cl_uint num_entries,  
                 cl_device_id *devices,  
                 cl_uint *num_devices)
```

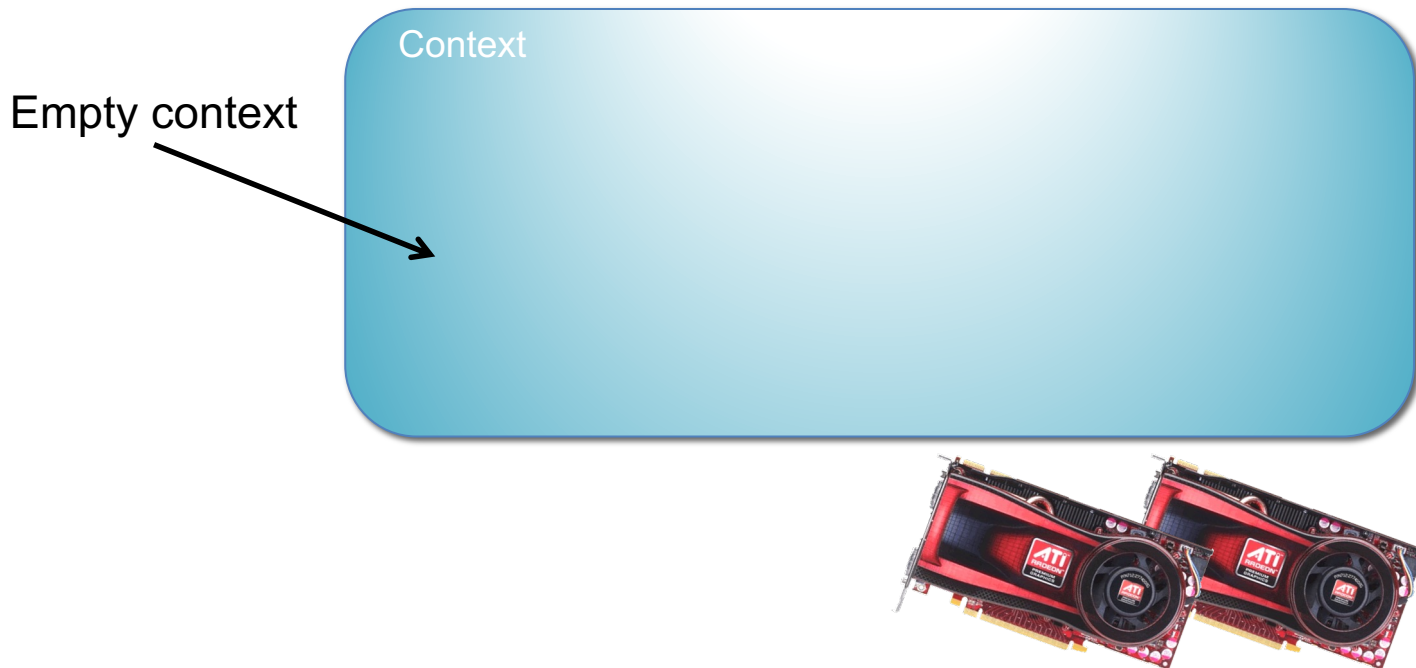
- We can specify which types of devices we are interested in (e.g. all devices, CPUs only, GPUs only)
- This call is performed twice as with clGetPlatformIDs
  - The first call is to determine the number of devices, the second retrieves the device objects

# Contexts

- A context refers to the environment for managing OpenCL objects and resources
- To manage OpenCL programs, the following are associated with a context
  - Devices: the things doing the execution
  - Program objects: the program source that implements the kernels
  - Kernels: functions that run on OpenCL devices
  - Memory objects: data that are operated on by the device
  - Command queues: mechanisms for interaction with the devices
    - Memory commands (data transfers)
    - Kernel execution
    - Synchronization

# Contexts

- When you create a context, you will provide a list of devices to associate with it
- For the rest of the OpenCL resources, you will associate them with the context as they are created



# Contexts

```
cl_context      clCreateContext (const cl_context_properties *properties,  
                                cl_uint num_devices,  
                                const cl_device_id *devices,  
                                void (CL_CALLBACK *pfn_notify)(const char *errinfo,  
                                                                const void *private_info, size_t cb,  
                                                                void *user_data),  
                                void *user_data,  
                                cl_int *errcode_ret)
```

- This function creates a context given a list of devices
- The properties argument specifies which platform to use (if NULL, the default chosen by the vendor will be used)
- The function also provides a callback mechanism for reporting errors to the user

# Command Queues

- A *command queue* is the mechanism for the host to request that an action be performed by the device
  - Perform a memory transfer, begin executing, etc.
- A separate command queue is required for each device
- Commands within the queue can be synchronous or asynchronous
- Commands can execute in-order or out-of-order



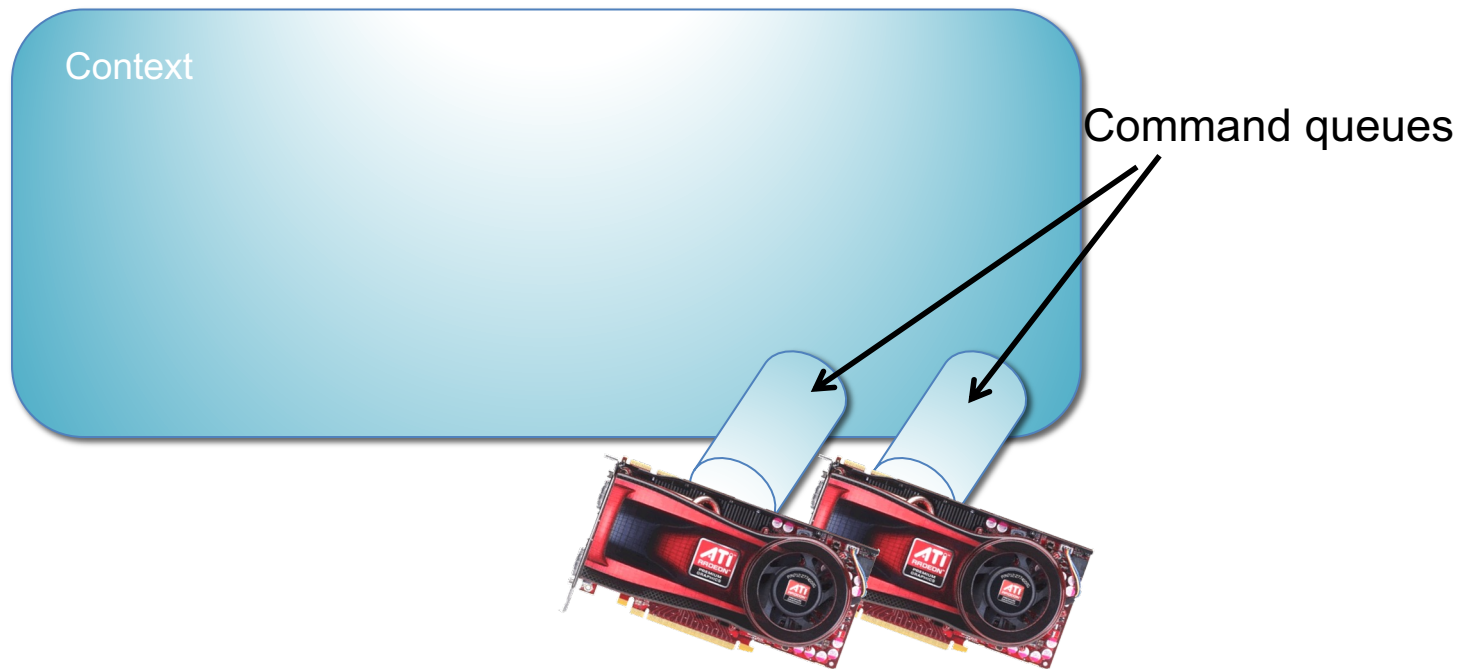
# Command Queues

```
cl_command_queue  clCreateCommandQueue (cl_context context,  
                                         cl_device_id device,  
                                         cl_command_queue_properties properties,  
                                         cl_int *errcode_ret)
```

- A command queue establishes a relationship between a context and a device
- The command queue properties specify:
  - If out-of-order execution of commands is allowed
  - If profiling is enabled

# Command Queues

- Command queues associate a context with a device
  - Despite the figure below, they are not a physical connection



# Memory Objects

- Memory objects are OpenCL data that can be moved on and off devices
  - Objects are classified as either buffers or images
- Buffers
  - Contiguous chunks of memory – stored sequentially and can be accessed directly (arrays, pointers, structs)
  - Read/write capable
- Images
  - Opaque objects (2D or 3D)
  - Can only be accessed via `read_image()` and `write_image()`
  - Can either be read or written in a kernel, but not both

# Creating buffers

```
cl_mem  clCreateBuffer (cl_context context,  
                        cl_mem_flags flags,  
                        size_t size,  
                        void *host_ptr,  
                        cl_int *errcode_ret)
```

- This function creates a buffer (cl\_mem object) for the given context
  - Images are more complex and will be covered in a later lecture
- The flags specify:
  - the combination of reading and writing allowed on the data
  - if the host pointer itself should be used to store the data
  - if the data should be copied from the host pointer

# Transferring Data

- OpenCL provides commands to transfer data to and from devices
  - `clEnqueue{Read|Write}{Buffer|Image}`
  - Copying from the host to a device is considered *writing*
  - Copying from a device to the host is *reading*
- The write command both initializes the memory object with data and places it on a device
  - The validity of memory objects that are present on multiple devices is undefined by the OpenCL spec (i.e. are vendor specific)
- OpenCL calls also exist to directly map part of a memory object to a host pointer

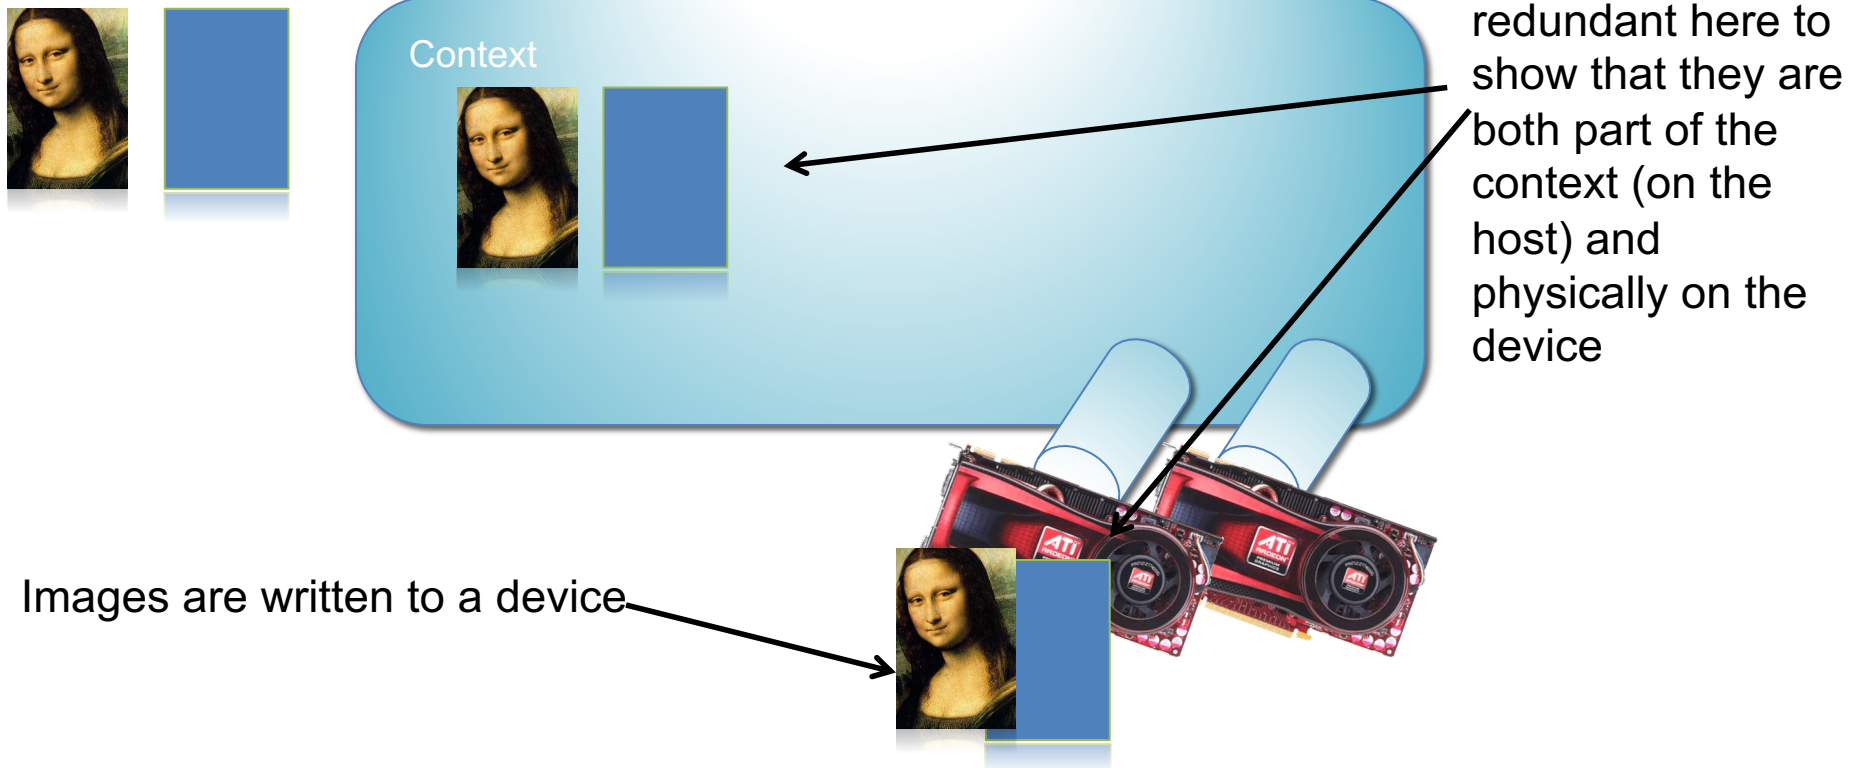
# Transferring Data

```
cl_int clEnqueueWriteBuffer (cl_command_queue command_queue,
                             cl_mem buffer,
                             cl_bool blocking_write,
                             size_t offset,
                             size_t cb,
                             const void *ptr,
                             cl_uint num_events_in_wait_list,
                             const cl_event *event_wait_list,
                             cl_event *event)
```

- This command initializes the OpenCL memory object and writes data to the device associated with the command queue
  - The command will write data from a host pointer (*ptr*) to the device
- The *blocking\_write* parameter specifies whether or not the command should return before the data transfer is complete
- Events can specify which commands should be completed before this one runs

# Transferring Data

- Memory objects are transferred to devices by specifying an action (read or write) and a command queue
- The validity of memory objects that are present on multiple devices is undefined by the OpenCL spec (i.e. is vendor specific)



# Outline

- OpenCL概述
- **OpenCL模型**
  - 平台模型
  - 执行模型
  - 内存模型
  - 编程模型
- **OpenCL程序设计**
  - 程序设计流程
  - **编程实例**



# OpenCL 实例：向量加法

- The “hello world” program of data parallel programming is a program to add two vectors

$$C[i] = A[i] + B[i] \text{ for } i=1 \text{ to } N$$

- For the OpenCL solution, there are two parts
  - Host code
  - Kernel code

# OpenCL实例：向量加法

Address space  
qualifier

kernel qualifier

Global thread  
index

Vector addition

```
// OpenCL Kernel Function for element by element vector addition
__kernel void VectorAdd(__global const float* a,
                        __global const float* b,
                        __global float* c,
                        int iNumElements)
{
    // get index into global data array
    int iGID = get_global_id(0);

    // bound check (equivalent to the limit on a 'for' loop for standard/serial C code)
    if (iGID >= iNumElements)
    {
        return;
    }

    // add the vector elements
    c[iGID] = a[iGID] + b[iGID];
}
```

# OpenCL实例：向量加法

Setup kernel grid

Allocate host resources

Create device context

Allocate device resources

Populate device memory

```
int main(int argc, char **argv) {  
    // set and log Global and Local work size dimensions  
    localWorkSize = 256;  
    globalWorkSize = RoundUp(localWorkSize, iNumElements);  
  
    // Allocate and initialize host arrays  
    srcA = malloc(sizeof(cl_float) * globalWorkSize);  
    srcB = malloc(sizeof(cl_float) * globalWorkSize);  
    dst = malloc(sizeof(cl_float) * globalWorkSize);  
    FillArray(srcA, iNumElements);  
    FillArray(srcB, iNumElements);  
  
    // Create the context  
    clGetPlatformIDs(1, &platform, NULL);  
    clGetDeviceIDs(platform, CL_DEVICE_TYPE_GPU, 1, &device, NULL);  
    context = clCreateContext(0, 1, &device, NULL, NULL, &ciErr1);  
  
    // Create a command-queue  
    cqCommandQueue = clCreateCommandQueue(context, device, 0, &ciErr1);  
  
    // Allocate the OpenCL buffer memory objects  
    dSrcA = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(cl_float) * globalWorkSize, ...);  
    dSrcB = clCreateBuffer(context, CL_MEM_READ_ONLY, sizeof(cl_float) * globalWorkSize, ...);  
    dDst = clCreateBuffer(context, CL_MEM_WRITE_ONLY, sizeof(cl_float) * globalWorkSize, ...);  
  
    // Copy data to GPU device  
    clEnqueueWriteBuffer(cqCommandQueue, dSrcA, CL_FALSE, 0, sizeof(cl_float) * globalWorkSize, srcA, ...);  
    clEnqueueWriteBuffer(cqCommandQueue, dSrcB, CL_FALSE, 0, sizeof(cl_float) * globalWorkSize, srcB, ...);  
}
```

# OpenCL实例：向量加法

Build kernel  
program

Set kernel  
arguments

Launch kernel  
execution

Read destination  
buffer

```
// Build the program
progSrc = oclLoadProgSource("VectorAdd.cl", ...);
prog = clCreateProgramWithSource(context, 1, progSrc, ...);
clBuildProgram(prog, 0, NULL, NULL, NULL, NULL);

// Create the kernel
kernel = clCreateKernel(prog, "VectorAdd", &ciErr1);

// Set kernel the Arguments
clSetKernelArg(kernel, 0, sizeof(cl_mem), (void*)&dSrcA);
clSetKernelArg(kernel, 1, sizeof(cl_mem), (void*)&dSrcB);
clSetKernelArg(kernel, 2, sizeof(cl_mem), (void*)&dDst);
clSetKernelArg(kernel, 3, sizeof(cl_int), (void*)&iNumElements);

// Launch kernel
clEnqueueNDRangeKernel(cqCommandQueue, kernel, 1, NULL, &globalWorkSize, &localWorkSize,...);

// Synchronous/blocking read of results, and check accumulated errors
clEnqueueReadBuffer(cqCommandQueue, dDst, ...);

// Cleanup
Cleanup();
return 0;
```

# 参考资料

- OpenCL异构并行计算：原理、机制与优化实践陈轶,吴长江,刘文志 著
- Khronos OpenCL Registr, <https://www.khronos.org/registry/OpenCL/>